

Retrieval Systems for Structured Data

The critical missing  for coupling LLM-powered query interfaces for factual data

Madelon Hulsebos

BIDS Seminar
October 2024

Finding the right data for basic questions or deep analysis is not easy.

1. How to help data analysts find their data for analytics tasks?
2. How to ground LLM-powered query interfaces in structured data?


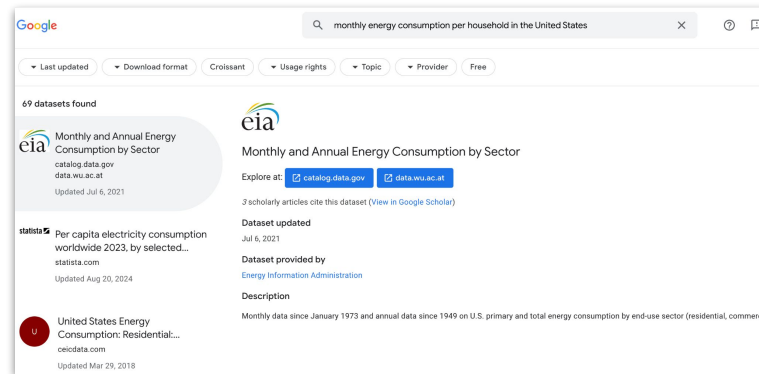
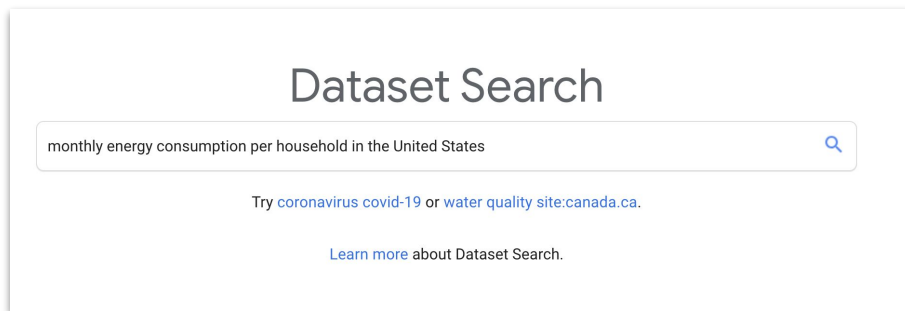
Part I:

How to help data analysts find their data for analytics tasks?

Current dataset search systems for analytics use-cases

Imagine you are looking for a dataset to forecast monthly energy consumption per household until May 2025.



Keyword Search and Metadata Filters



69 datasets... Which one?


Why is dataset search still so hard in practice?

89 data practitioners recruited through social media & mailing lists



How do Practitioners Search?


How do you search?



“...Also, identify **people** who have worked on **similar problems**...”

“... I ask more **experienced colleagues** which ones are most **inherent** to the analysis...”

Key challenges with existing systems?



“The biggest challenge I’ve noticed is **messy variable naming**...”

“**Categorical level of detailing** is required...”

“There are **too many table results** after the initial search....”

“**Not many features to search/query keywords**...”

Towards Next-Generation Dataset Search Systems

Task-driven: explicit **data needs unknown to user**, start from the task

Hybrid: search spans **diff representations** of a table; raw metadata + **semantics**

Iterative: beyond keywords+filters; complex process, need for more pruning

Comprehensibility+diversity: result sets hard to **digest and navigate; IR principles?**

Task-driven Search with Hypothetical Schema Embeddings

(1) Task-driven query

What data is needed to **train a machine learning model** to **forecast demand for medicines across suppliers**?

(2) Hypothetical Schema generation

Instruction: generate schema needed for the given task query.

Query: {task-driven query}

LLM output: "hypothetical_schema"

medication table: medication id, medication name, ...

sales table: medication id, supplier id, date, quantity sold, ...

(3) Embed(hypothetical_schema)

(4) Retrieve source tables from vector store similar to hypothetical_schema

Iterative Search through Query Breakdown

Initial complex search query

A dataset to train a **medicine forecast model**, should **contain a, b, c**, and **span 2 years** with a **date range e to b**.
The fact table should be at **r granularity** and contain **20,000 records**.

“A dataset for **task x**,

<retrieved tables>

Retrieve through Hypothetical Schema Embeddings

Data should contain **a, b, c**

<pruned tables>

prune: schema similarity

Data should span **2 years** with a **date range e to b**.

<pruned tables>

prune: metadata **text-to-sql**

The fact table should be at **r granularity** and **contain 20,000 records**.”

<pruned tables>

prune: metadata **text-to-sql**

LLM-Assisted Interface for Dataset Search

Things we want from Dataset Search interfaces:

D1 LLM Elicitation through Proactive Guidance

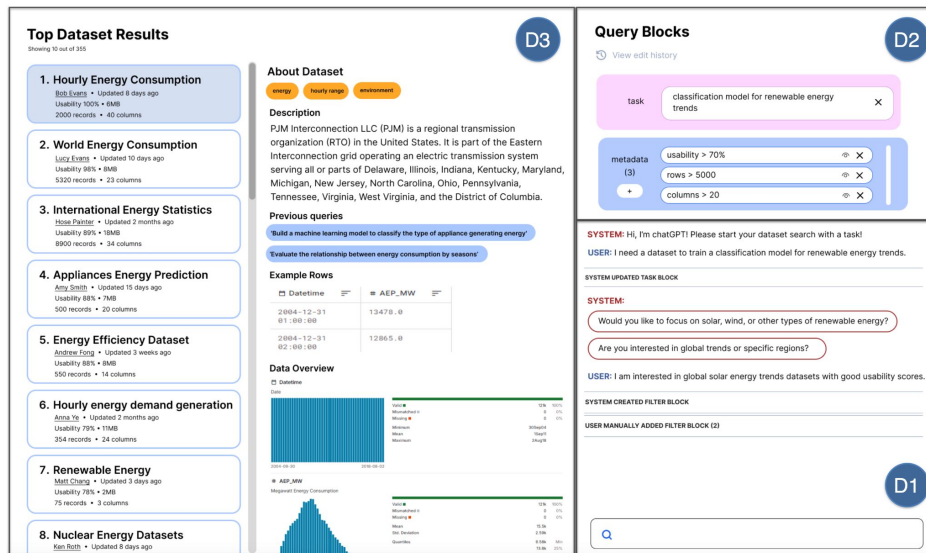
Purpose: Prompt users to share more information about their needs, which will be reflected in the query blocks & search interface.

D2 Dynamic Query Decomposition

Purpose: Allow users to see how the LLM is dynamically updating and refining the search space, providing transparency into the search process.

D3 Allowing Users to Compare Datasets Efficiently

Purpose: Facilitate high-level exploration of datasets by organizing them into topics and enable users to delve into metadata details of individual datasets as they iteratively build and refine their queries.



Part 2:

How to ground LLM-powered query interfaces in structured data?

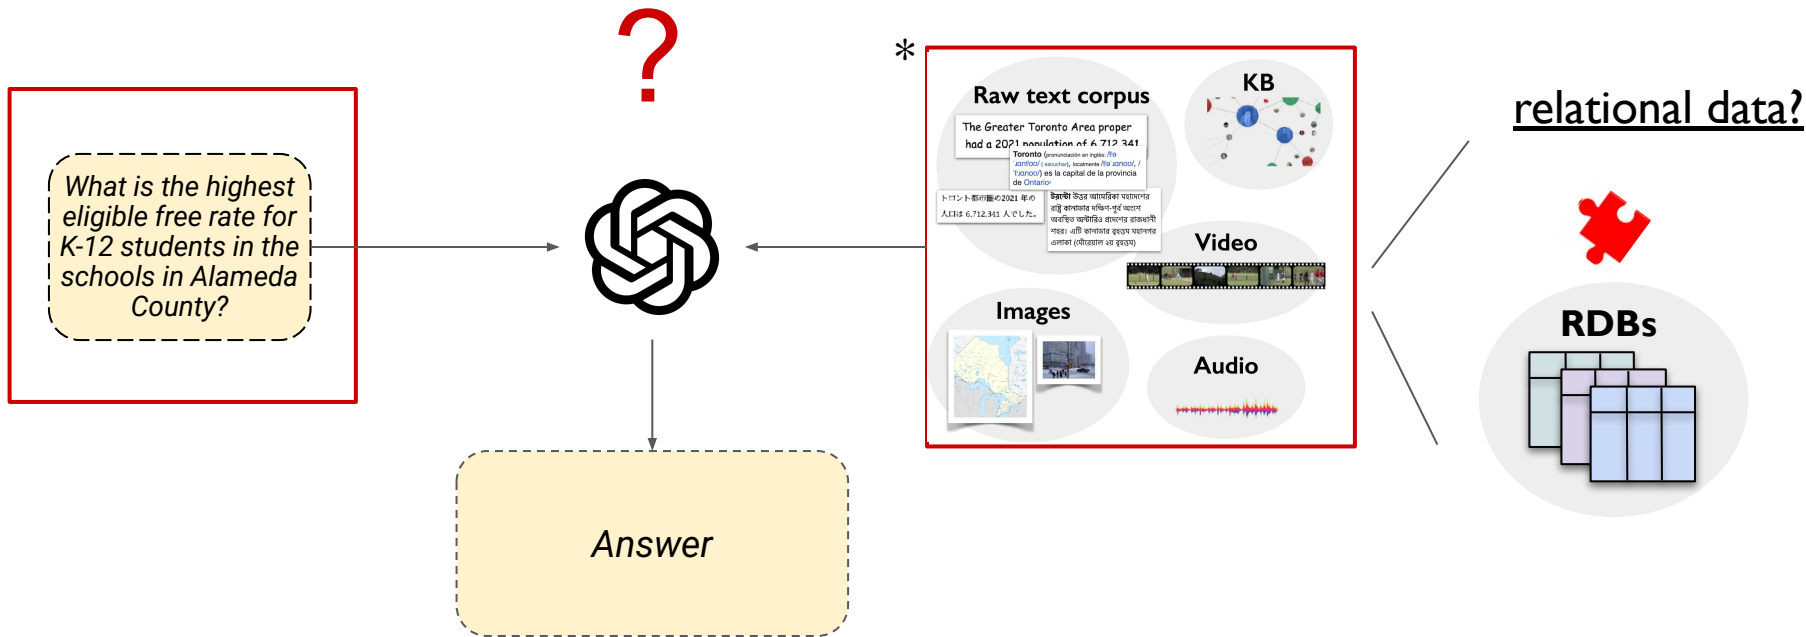
Asking LLMs complex questions

What is the highest eligible free rate for K-12 students in the schools in Alameda County?



*".... To determine the highest free rate specifically in Alameda County schools, **you'd generally need data from specific school districts or schools in the area**, as this rate can vary widely depending on the socio-economic demographics of each district. ..."* *

We need “specific” data to ground LLMs



Why we need RAG over *structured data*

A pattern in **practice**: “**everyone** cares about structured data”.

Structured data serve **high-value** insights! Up-to-date, domain-specific, facts...

Queries & RAG pipeline

“Which **urban Japanese prefecture** is not associated with **thorny trees**?” [table lookup]

“Shane Hall ran a total of **190** races between the year of **1995 - 2008**” [aggregate & compare]

“What is the **highest eligible free rate** for K-12 students in the schools in **Alameda County**” [aggregate]



Retrieval is difficult, but crucial...

“.. keep in mind that a good RAG system is really **hard to build**.


If your **retrieval system is mediocre**,

the **retrieval can easily distract LLMs** to backfire...

There is still a long way to go.” - Wenhui Chen (Univ of Waterloo)

Important grounds to explore...

Retrieval/generation complexity depends on query



- What “task” does the **query intend** to do?
- How should we **process** the table(s), relational DBs?
- **What should we embed** of the table(s) and metadata, and **how**?
- Given query and embedded corpus, **how to retrieve relevant table**?
- **Which data source** to retrieve from, and when?
- (How) should methods, models, systems **generalize across tasks, datasets**?

Methods for table retrieval

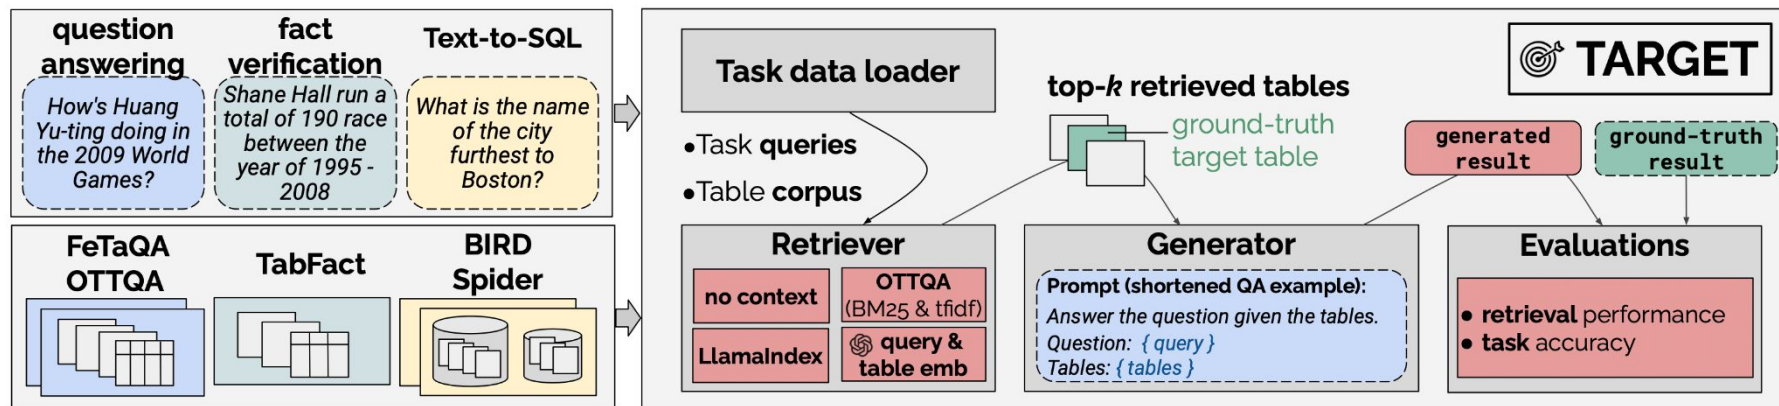
① Embedding of tables in corpus, and input query

- BM25 / TF-IDF (sparse lexical representations)
- Generate summary/metadata → embed summary + table
- “Naive” embedding of table (header / header+rows) and query

② Similarity search (e.g. cosine similarity) to identify top- k relevant tables

But how effective are these? How robust across datasets and tasks? No one really knows!

TARGET: Benchmarking Table Retrieval for Generative Tasks



- Diverse: tasks & datasets
- Extensible: easily add new tasks, new datasets
- Adaptable: eval custom retriever, generator

TARGET insights

Method	Question Answering						Fact Verification			Text-to-SQL					
	OTTQA			FeTaQA			TabFact			Spider		EX		BIRD	
	R@10	s	SB	R@10	s	SB	R@10	s	P/R/F1	R@1	s	EX	R@1	s	EX
No context	-	-	0.414	-	-	12.495	-	-	0.578/0.42/0.44	-	-	0	-	-	0
OTT-QA BM25	0.955	0.001	0.606	0.082	0.001	1.631	0.338	0.001	0.75/0.26/0.39	0.635	0.001	0.385	0.709	0.001	0.181
w/o table title	0.443	0.001	0.529	0.084	0.001	1.555	0.331	0.001	0.75/0.26/0.38	0.5	0.001	0.376	0.535	0.001	0.164
OTT-QA TF-IDF	0.950	0.001	0.425	0.083	0.001	1.639	0.336	0.001	0.75/0.26/0.38	0.622	0.001	0.474	0.640	0.001	0.227
w/o table title	0.43	0.001	0.593	0.083	0.001	1.527	0.322	0.001	0.75/0.25/0.37	0.492	0.001	0.376	0.491	0.001	0.164
LlamaIndex	0.458	0.354	0.507	0.435	0.396	13.745	0.827	0.297	0.73/0.34/0.47	0.735	0.198	0.559	0.937	0.228	0.311
OpenAI embedding	0.950	0.190	0.599	0.722	0.200	17.64	0.779	0.189	0.76/0.51/0.61	0.768	0.193	0.602	0.926	0.199	0.317
header only	0.950	0.189	0.61	0.718	0.18	17.66	0.781	0.187	0.75/0.48/0.58	0.833	0.175	0.646	0.958	0.191	0.323

- BM25/TF-IDF **less effective**, only with very descriptive table name.
- Table rows can **“distract” embeddings**, particularly in RDBs as seen in practice.
- Generating summary/metadata can help, but **not all tables easy to LLM-summarize**.

Still much to explore...

- What is right input of (meta)data to not “distract” embedding?
- How do we route to proper data source, interpret the task, etc?
- **The reality in practice is much harder:**
 - How do methods perform on more *challenging tasks & datasets*?
 - Closing semantic gap $e(\text{query})$ and $e(\text{table})$; most public datasets relatively “easy” match between query and tables.
 - Relational databases are large → in-DB schema and table retrieval.

Roadmap for TARGET


TARGET is out **TODAY!**


RAG tables over tables with TARGET!

```
from target_benchmark.retrievers import AbsCustomEmbeddingRetriever
class YourRetriever(AbsCustomEmbeddingRetriever):
    def __init__(self, **kwargs):
        # load your favorite table retriever!


    def retrieve(self, query: str, dataset_name: str, top_k: int):
        # given a query, retrieve the top-k table id

    def embed_corpus(self, dataset_name: str, corpus: Iterable[Dict]):
        # use retriever, embedding models, etc. to embed the corpus!
```

- Ready to eval table retrieval and e2e generation: **input welcome for v2**
- Data on HF, code on GH,  `pip install target_benchmark`
- <https://target-benchmark.github.io>



Key takeaways

- Retrieval (RAG, agents, or dataset discovery) is critical to get insights from data.
- Dataset discovery requires co-worker consultation. Existing interfaces are inflexible. We propose lightweight retrieval w **Hypothetical Schema Embeddings**, and a **flexible LLM-assisted interface** that takes a *task* and *doesn't fix metadata filters*.
- To push table retrieval for grounding LLMs in structured data forward, we introduce  **TARGET: the first benchmark for RAG for structured data**, and uncover deficits of methods that are robust for unstructured text! More to come...